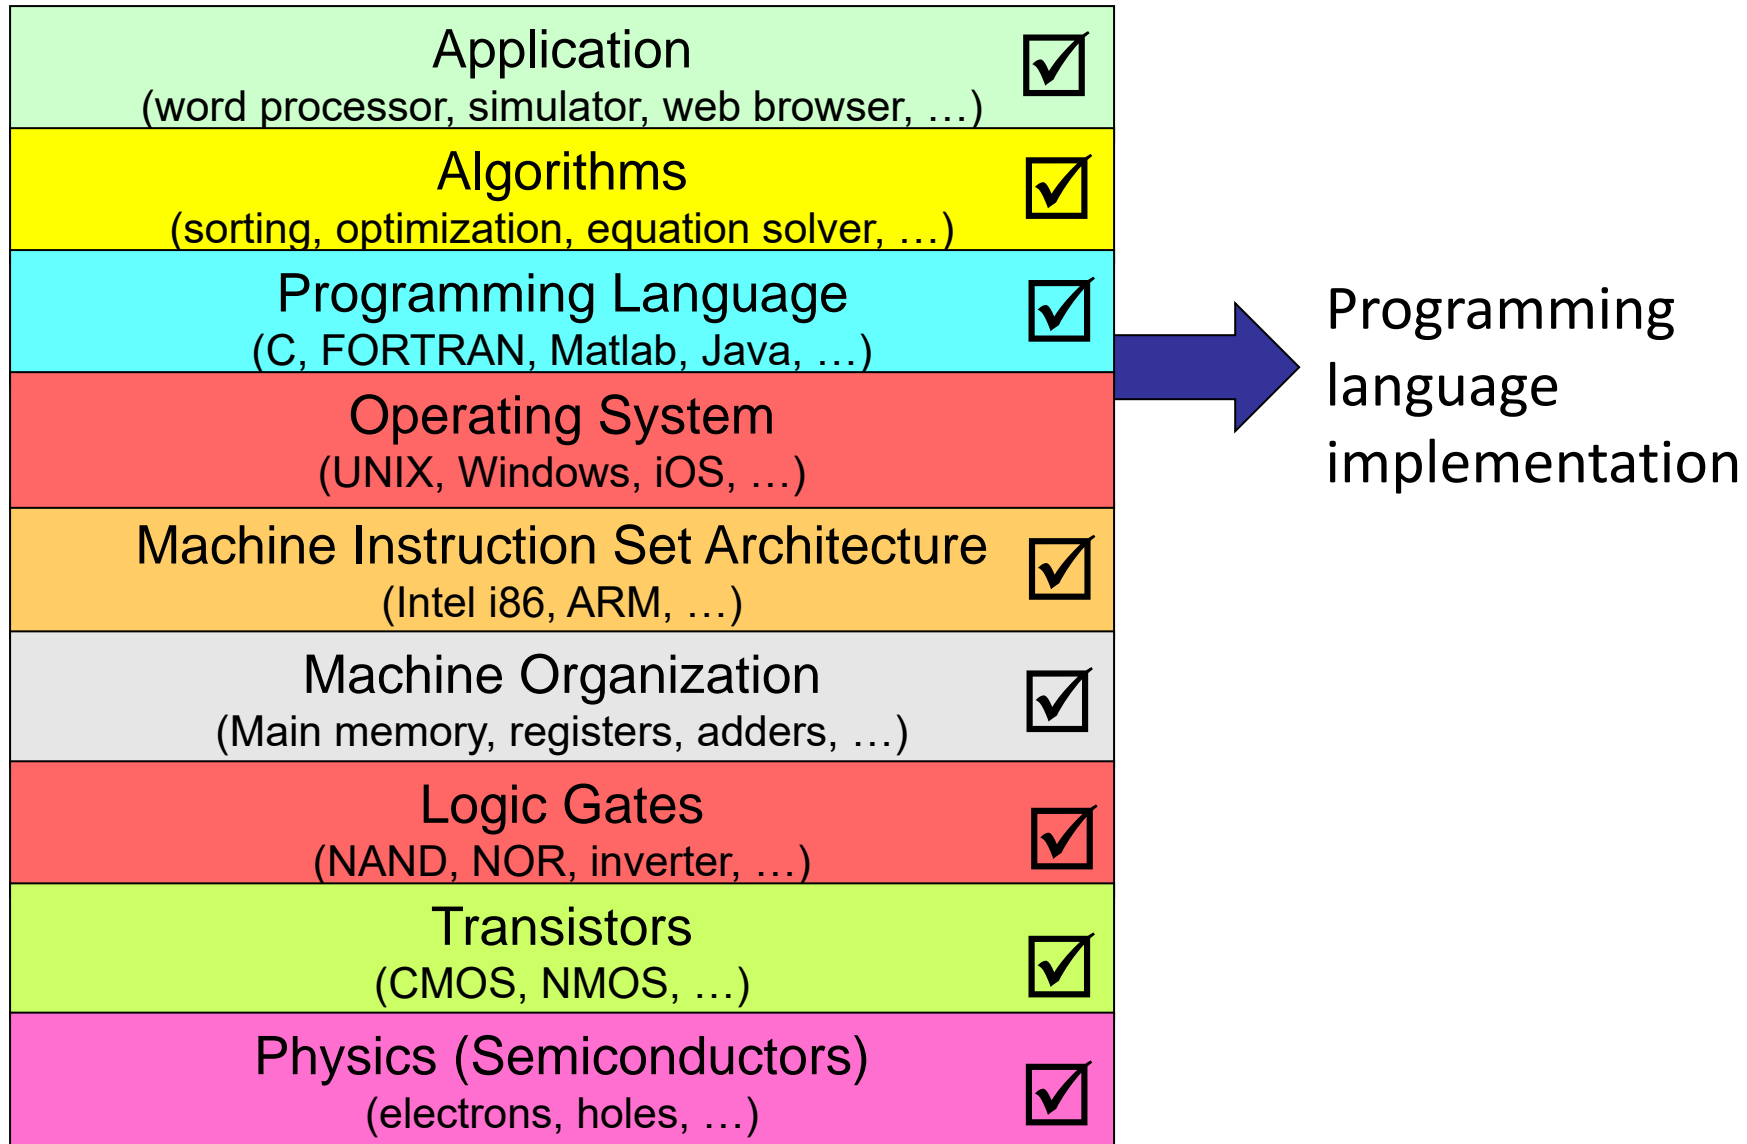


# Functions and I/O


For use in CSE6010 only

Not for distribution

# Levels of Abstraction in Computers



# Memory Map



- Layout similar to that typically used in Unix systems
- Runtime stack and heap grow/shrink during program execution

\* Not related to heap data structure

# Structure of a Typical C Program

```
int a, b, c;                /* global variables */  
// function prototypes go here
```

```
main ()  
{   int i, j, k;  
    i=1; j=2; k=3;  
    . . . foo(i, &j) . . .      /* function call (invocation) */  
}
```

arguments: actual values passed

parameters: variables in func definition

```
int foo(int x, int *y) /* function definition */  
{   int a;             /* local variable */  
    a=4;                /* reference local variable */  
    . . . bar(a); . . .  
    return (a);  
}
```

```
void bar (int w)  
{   double i;  
    . . . reference a . . .      /* reference to global*/  
}
```

# Parameter Passing: Call By Value

Call by Value (used in C)

- A *copy* of the argument is created, passed to function
- Upon return, the copy is discarded
- Implications of copying
  - Copying can be expensive (time and space) for large parameters
  - Modifications to the parameter within the function are made to the copy, and are **not** visible to the calling program
- Q: What if you want modifications to be visible to the caller?
- A: Pass a *pointer* to the data being passed; value passed is an address, e.g., `scanf ("%d", &i) ;`
- Aside: in C, array arguments pass a pointer to the array

Other approaches: call by value return, call by reference, call by name

# Call by Value Example

```
#include <stdio.h>
int foo (int x, int *y); // integer, pointer to integer

main()
{   int i, j, k;

    i=1; j=2; k=3;
    k = foo (i, &j);
    printf ("i=%d, j=%d\n", i, j);
}

int foo (int x, int *y)
{
    x=3;
    *y = 4;
    return (...)
}
```

What is printed?

i=1, j=4

# Calling a Function

- What is needed to properly implement a function call and the return to the main program?

# Calling a Function

- What is needed to properly implement a function call and the return to the main program?
  - Keeping track of where we are in the program: need to know the address where the function begins, the address of the next instruction in the main program (incremented PC) for returning, and the values of local variables in the main program
  - Variable space for the function: track arguments passed, variables returned; know and allocate space for local variables used in the function
- We will be talking about both of these.



# Implementing Function Call / Return


```
main()  
{  
    k = foo (i, &j);  
}  
  
int foo (int x, int *y)  
{  
    return (...)  
}
```

- Function Call
  - Similar to branch instruction
  - Execution “jumps” to the code implementing the function
- Function Return
  - Where does execution continue after the function is completed?
- Implementation solution
  - Function call saves the **return address** (where to resume)
  - Jump to the saved return address to return from the function


# LC-3 Machine Instructions

## LC-3: save the return address in register R7

```
JSR      offset    // Jump Subroutine (=function) instruction  
           // R7 <- PC+1 (return address)  
opcode   // PC <- PC+1 + offset (sign extended)
```



```
RET          // PC <- R7
             // special case of the JMP instruction (R7)
```



# Function Call and Return

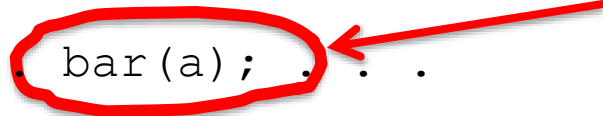
```
main ()  
{  
    . . . foo(i, &j) . . .  
}
```

JSR foo // Save return address (PC+1) in R7




```
int foo (int x, int *y)  
{  
    . . . bar(a); . . .  
}
```

Where do we save the return address for bar?



RET // Resume execution of caller (jump to address in R7)

```
void bar (int w)  
{  
    . . .  
}
```



# The Runtime Stack

- We can consider memory allocation for the return address for function calls like list data structure inserts and deletes
  - Function call: insert return address into a list
  - Function return: delete return address from list
- When a function returns, **the most recently saved return address** should be used
- This suggest a last-in-first-out approach, i.e., a **stack**

Storage for function calls is managed in an area of memory called the **runtime stack**.

# LC-3 Function Call / Return

```
main ()
```

```
{  
  . . . foo(i, &j) . . .  
}
```

JSR foo (save return address in R7)

Resume execution here

PUSH R7

```
int foo (int x, int *y)  
{
```

```
  . . .  
  bar(a);
```

JSR bar (save return address in R7)

Resume execution here

```
  . . .  
}
```

R7<-POP

RET (PC<-R7)

```
void bar (int w)  
{
```


PUSH R7

```
  . . .  
}
```

R7<-POP

RET (PC<-R7)

Runtime Stack



LC-3: Register R6 used as  
pointer to top of stack

# We're only just getting started!

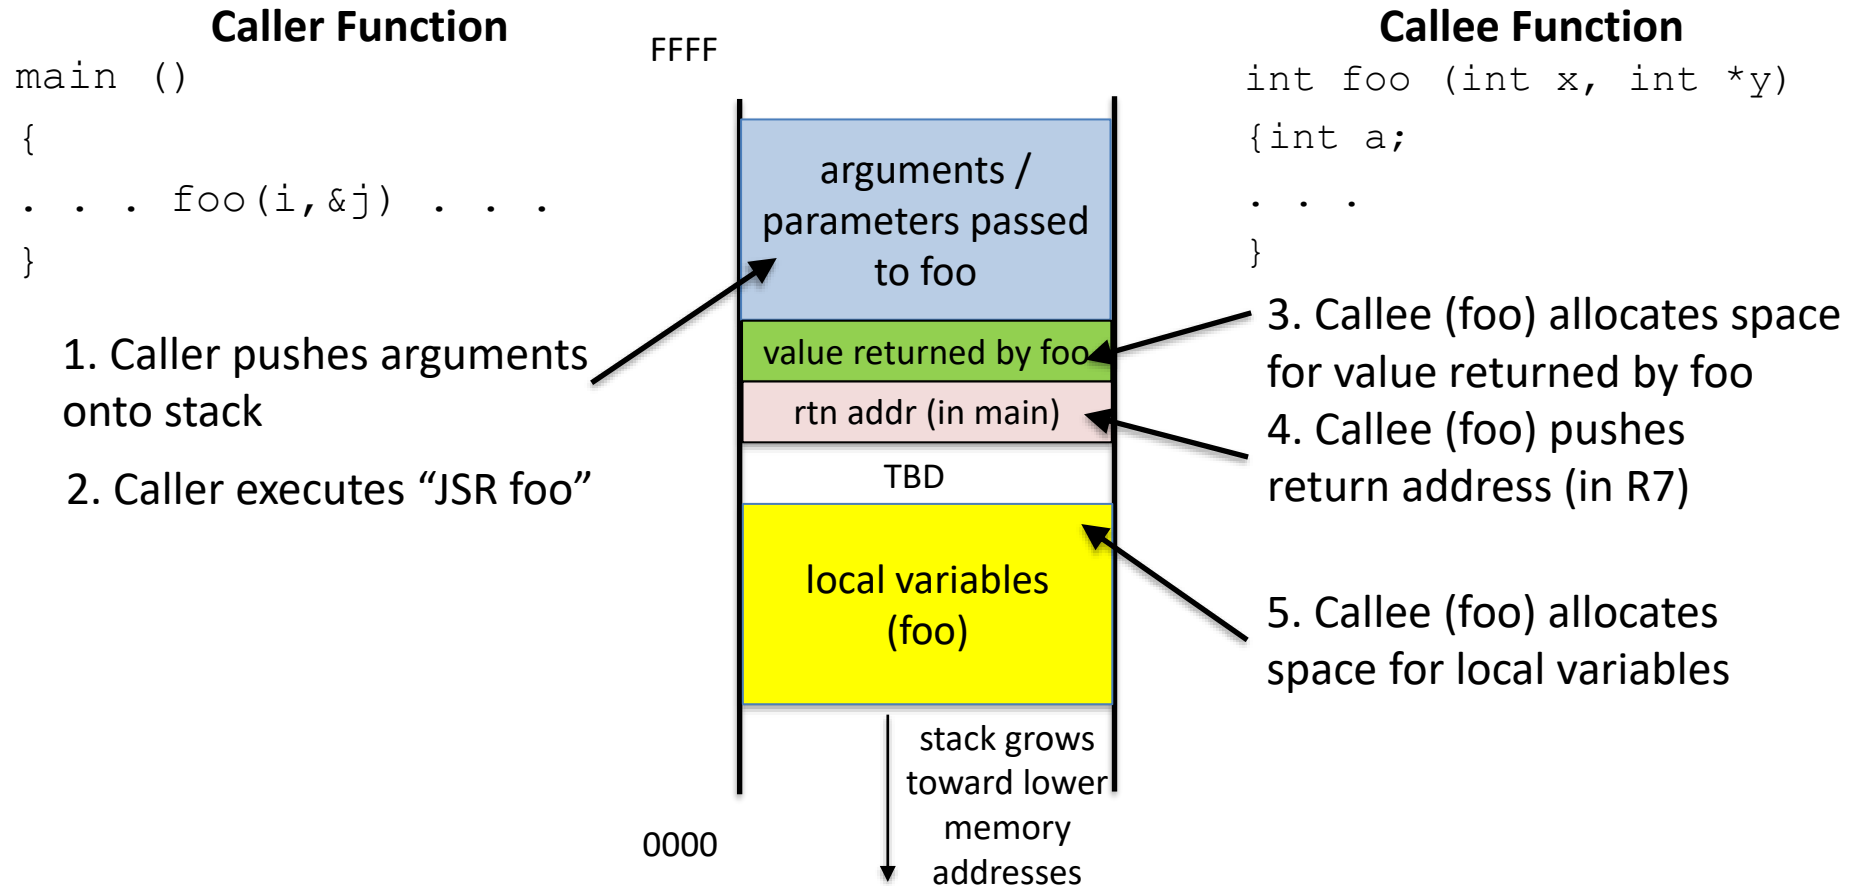
## Functions need additional storage.

Other storage must be allocated with each function call, and deallocated when the function returns

- Parameters/arguments passed to the function
  - Caller copies argument values into this storage
  - The callee (function being called) needs to access parameters
- Local variables declared within the called function
  - Allocated with each function call, released upon return
  - A function can call itself (recursion)
- Value returned by the function
  - Callee writes return value into this storage
  - Caller must be able to access the return value

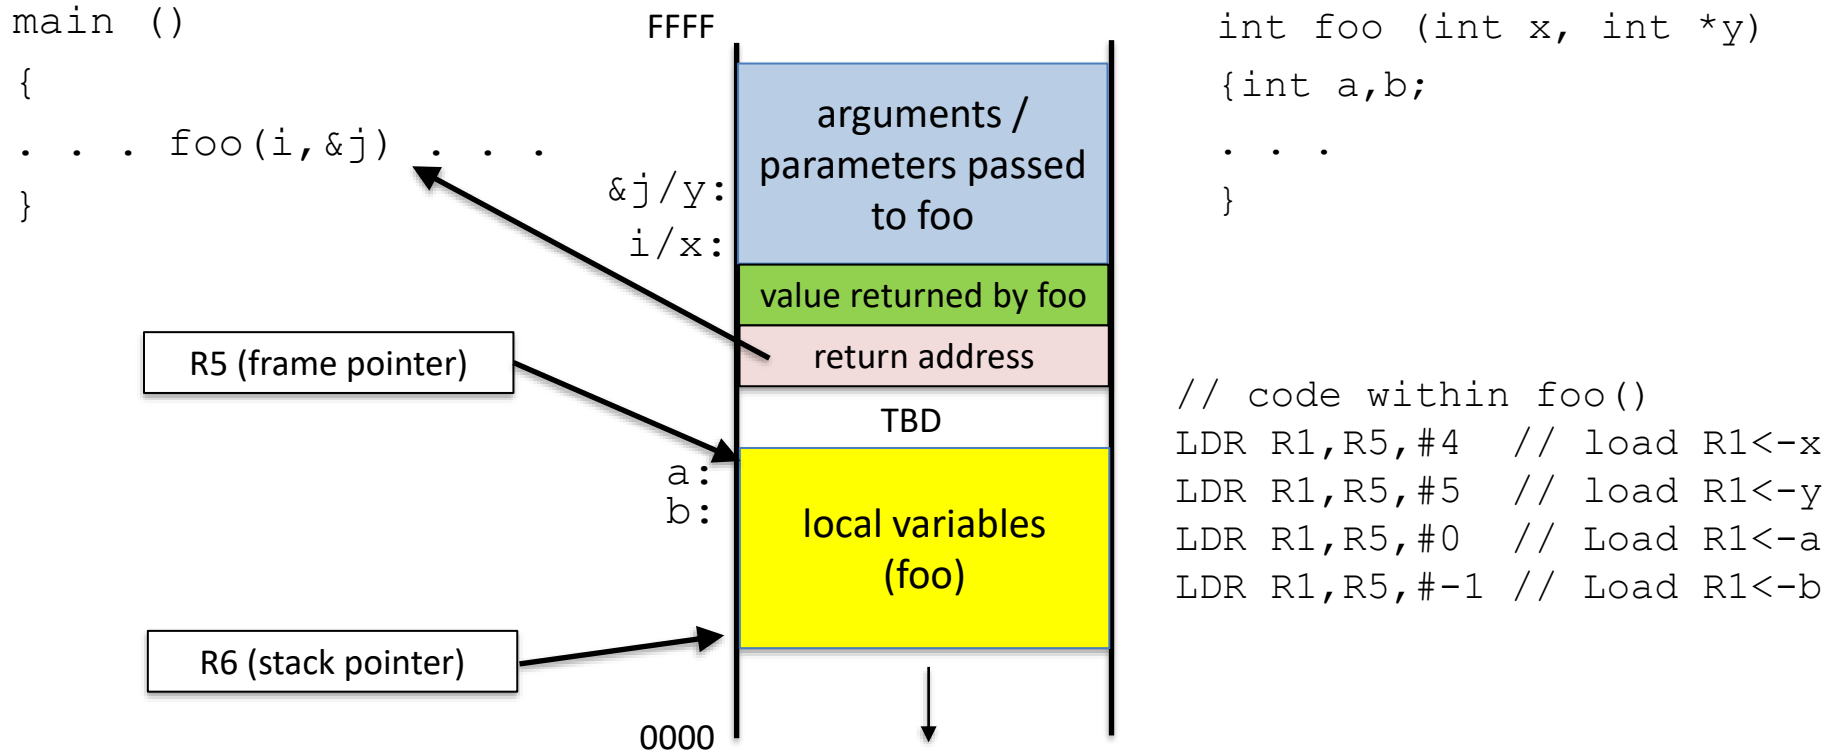
All of these items are also stored on the runtime stack

# Stack Frame (Activation Record)



- Holds parameters/arguments, local variables, return value, return address for an invocation (call) of a function
- One frame associated with each execution (invocation) of a function
  - At any instant, there is a single "current" frame being used
  - New frame pushed onto runtime stack with each function call
  - Pop frame from stack with each function return

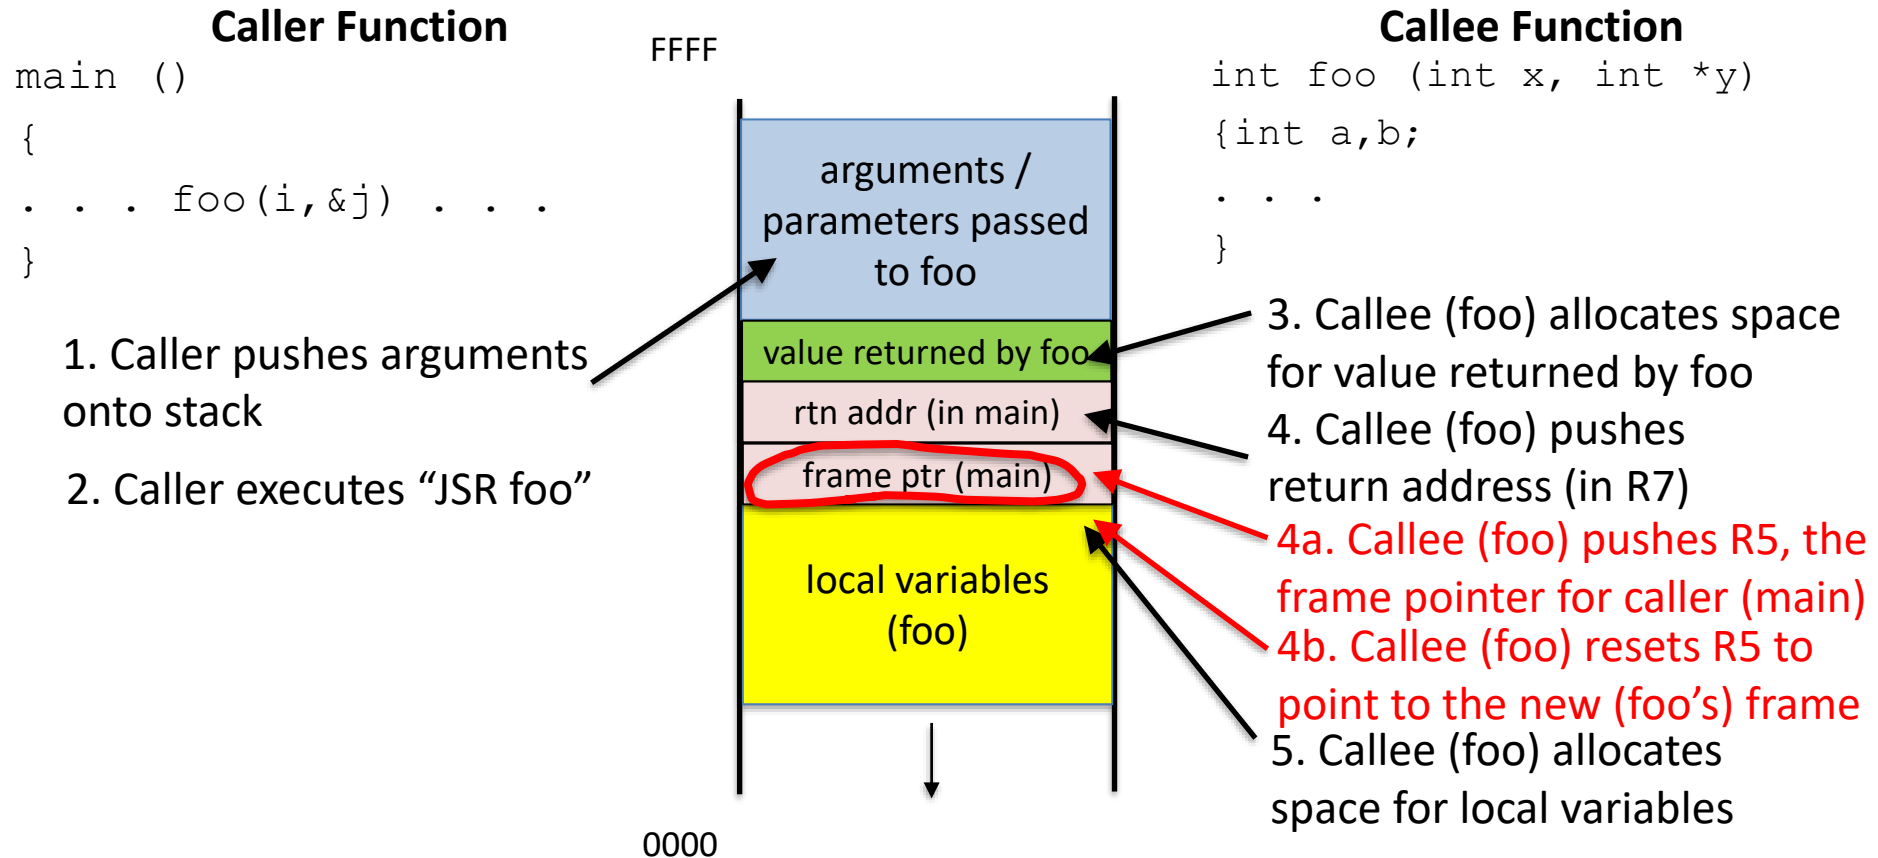
# Frame Pointer Register



- Each function invocation needs a way to access its parameters, local variables, return value (i.e., its stack frame)
- A CPU register (**frame pointer register**) holds a pointer to the current stack frame
  - LC-3 uses register R5
  - R5 points to first local variable in stack frame
  - Use indexed addressing mode to access locals, parameters, return value (sound familiar?)



# Stack Frame (Revised)



- The frame pointer points to the stack frame for the function **now** being executed
- Function call: need to save the frame pointer for the caller on the stack (in addition to the return address)

# Function Call Example

```
main()  
{  int i, j, k;
```

```
...
```

```
  i=1; j=2; k=3;
```

```
  k = foo (i, &j);
```

```
...
```

```
}
```

```
int foo (int x, int *y)
```

```
{
```

```
  int a;
```

```
...
```

```
  x=3; *y=4;
```

```
  a=4;
```

```
  bar (a);
```

```
...
```

```
  return (a);
```

```
}
```

```
void bar (int w)
```

```
{
```

```
...
```

```
}
```

*Call (invoke) function foo():*

- Push &j onto stack
- Push i onto stack
- JSR foo (save PC+1 in R7)

*Execute function foo():*

- **Allocate storage** on stack to hold the value returned by the function
- Push R7 (return address)
- Push R5 (frame pointer for main)
- Set R5 to point to new stack frame
- **Allocate storage** for local variables (a)
- Execute body of function (access local variables, parameters, etc.)

*Return:*

- Copy return value into return value storage
- Pop (**release storage**) for local variables
- Pop (restore) frame pointer into R5
- Pop return address into R7
- RET (jump to address in R7)

*Resume execution after call:*

- Pop (**release storage**) return value
- Pop (**release storage**) for parameters

# Runtime Stack Example

```
main()
{  int i, j, k;
```

```
...
    i=1; j=2; k=3;
    k = foo (i, &j);
```

```
...
}
    After foo call, i is still 1;
    j has been changed to 4
```

```
int foo (int x, int *y)
{
```

```
    int a; ...
    a=4;
    x=3;
    *y=4;
    bar (a);
```

```
...
}

void bar (int w)
{
...
}
```

Push new stack  
frame when foo  
is called

Pop stack frame  
when foo  
returns

Push new stack  
frame when bar  
is called

Pop stack frame  
when bar  
returns

	rest of main's stack frame; system space
FDEF (i)	1
FDEE (j)	<del>2</del> 4
FDED (k)	3
FDEC (y)	FDEE (addr of j)
FDEB (x)	<del>1</del> 3
FDEA	return value (foo)
FDE9	addr instr after JSR foo
FDE8	FDED (frame ptr, main)
FDE7 (a)	4
FDE6 (w)	4
FDE5	return value (bar)
FDE4	addr instr after JSR bar
FDE3	FDE7 (frame ptr, foo)

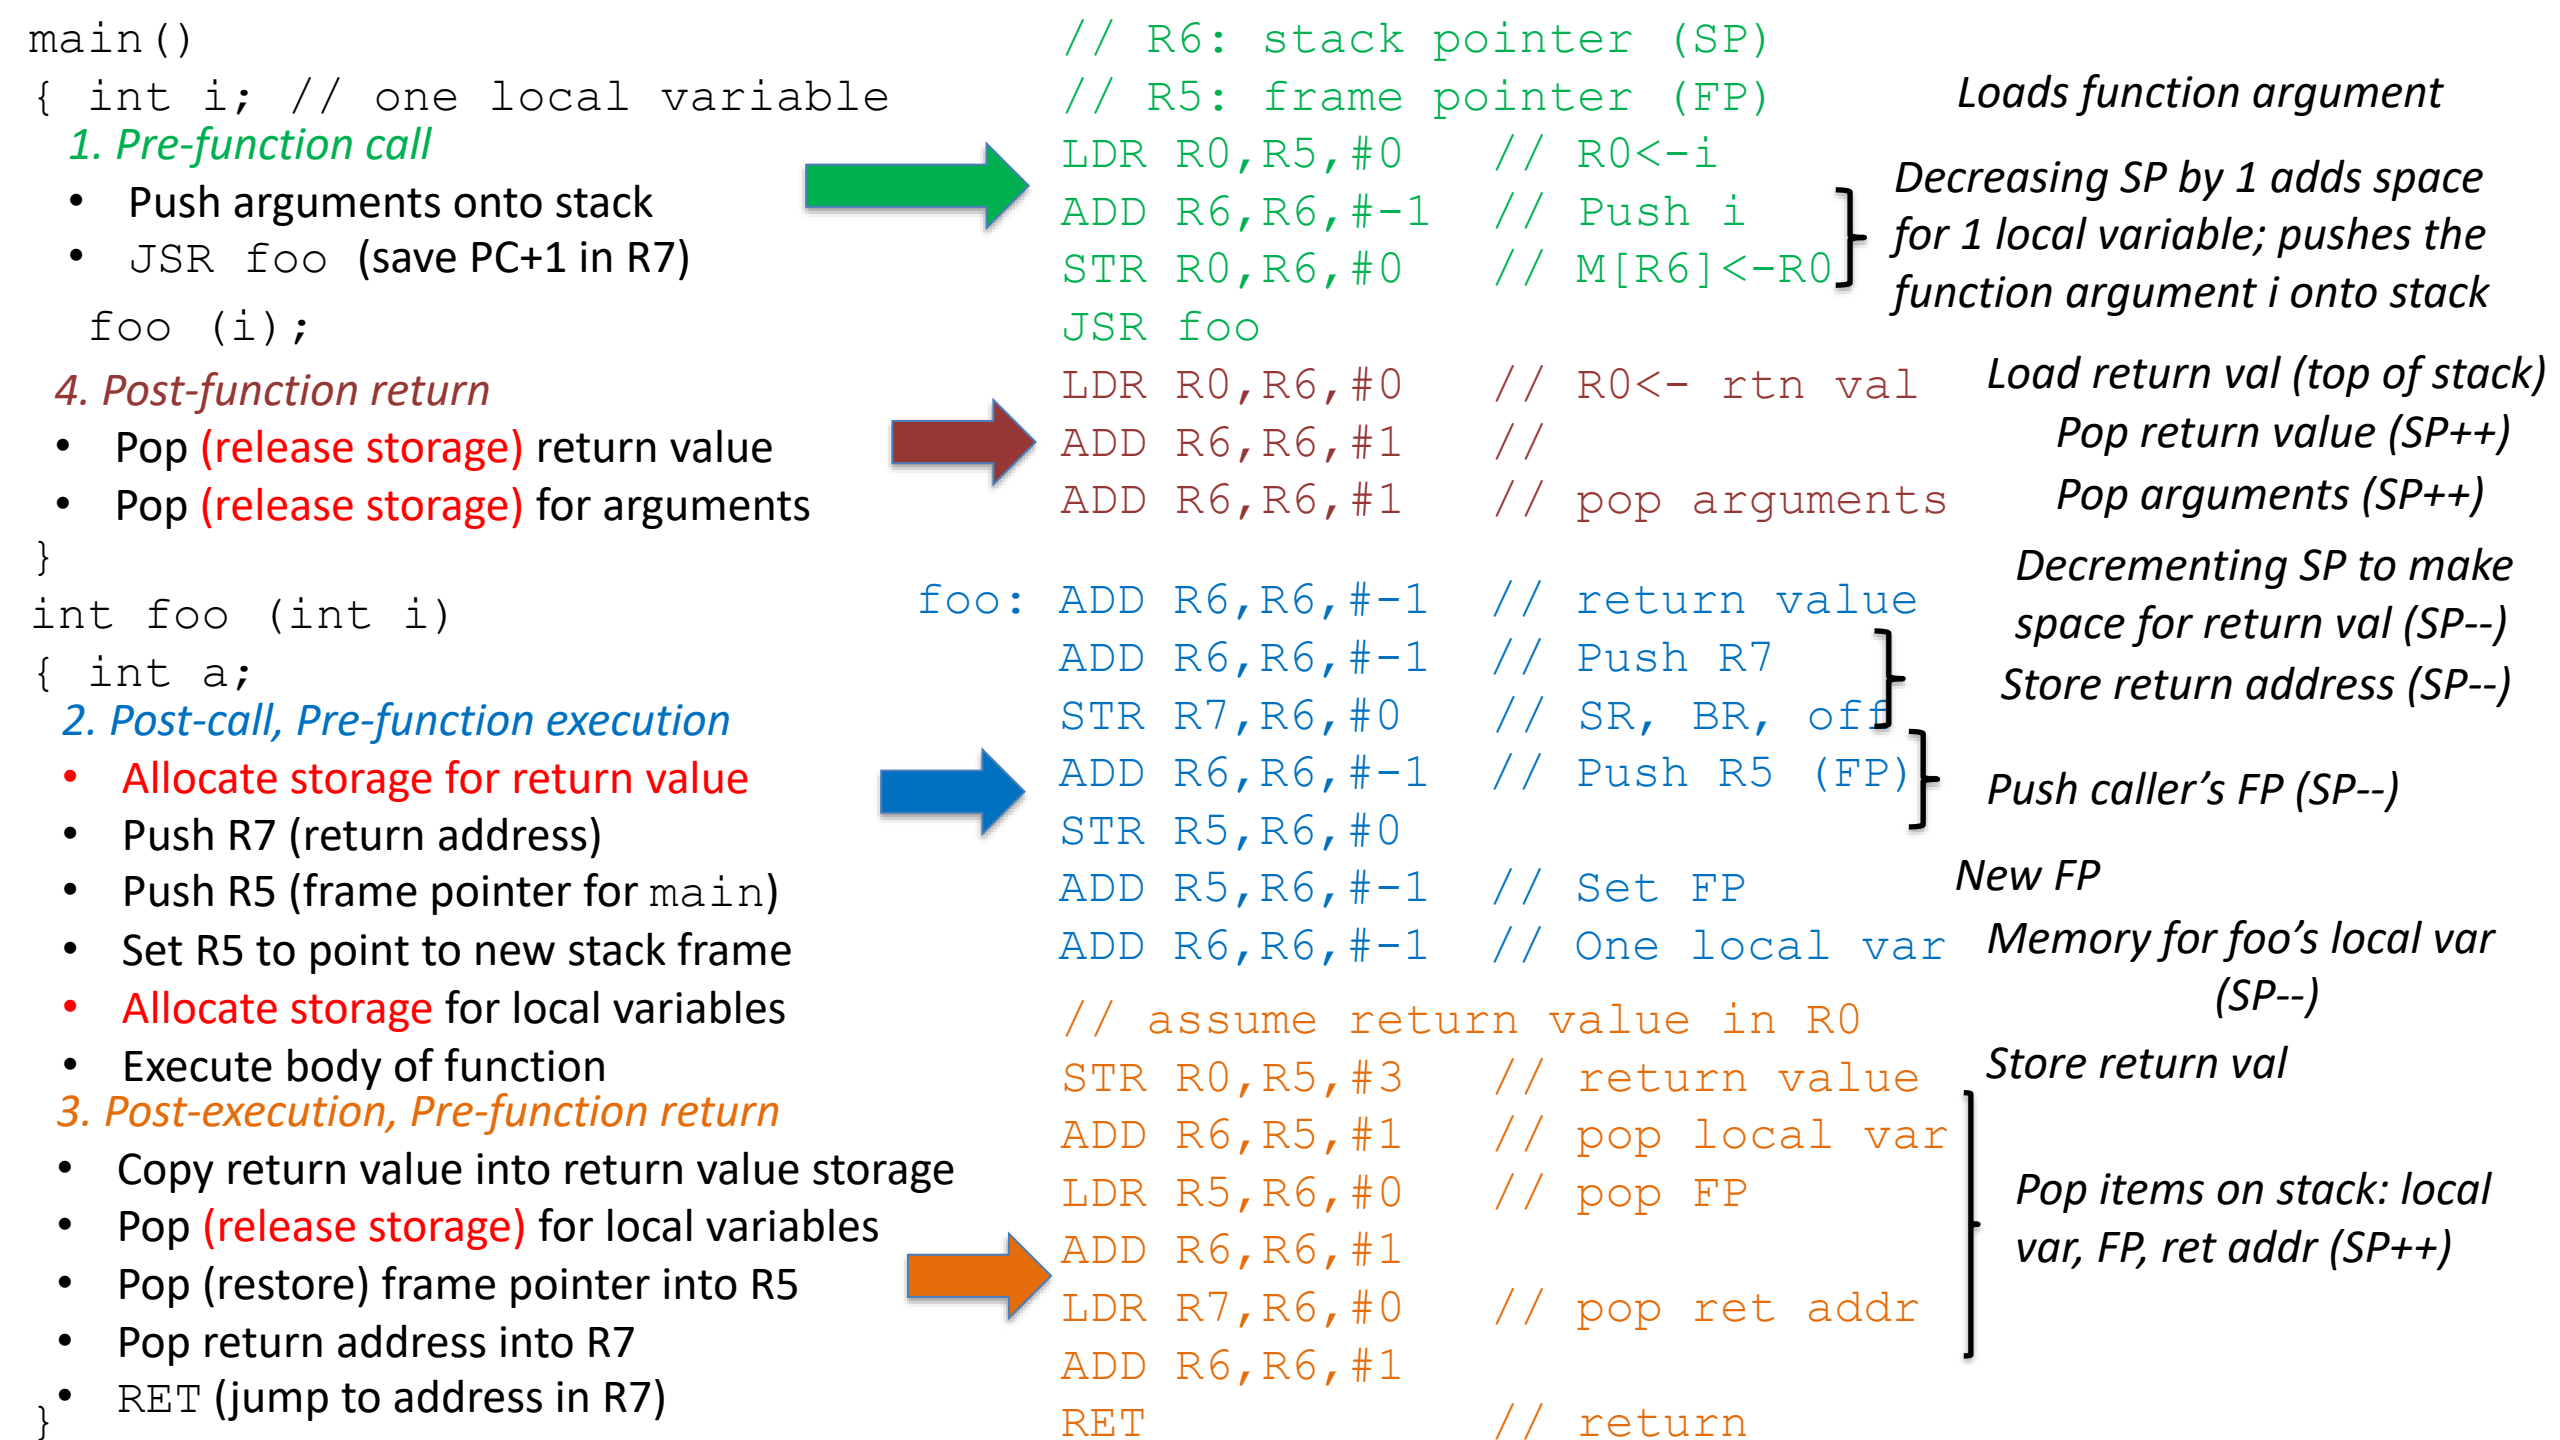
Stack Frame for  
main()

Stack Frame  
for foo()

Stack Frame  
for bar()

Note that  
addresses *within an  
activation record and  
on the runtime stack*  
usually **decrease**


- Function call: push new frame onto stack; this becomes the current frame
- Return: Pop current frame from stack



# Summary: Functions

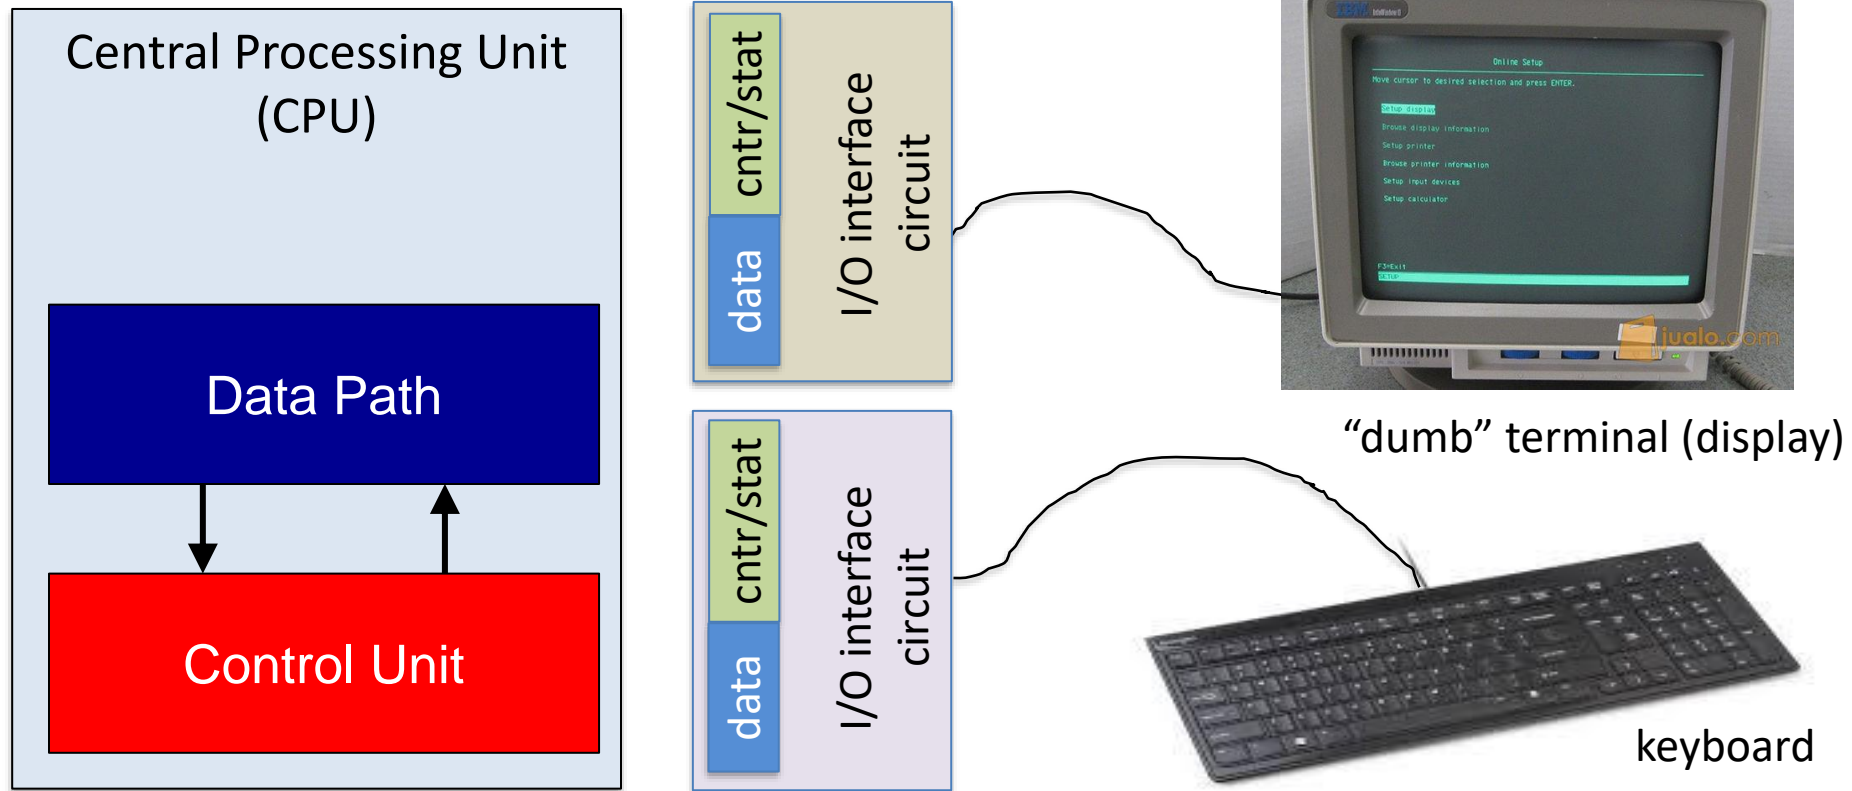
- C uses call by value; be careful whether to pass a variable or a pointer
  - If argument is to be modified by function, pass a pointer
  - If the argument is large (e.g., a large structure), pass a pointer
  - Otherwise, use the variable itself as the argument
- C (and many other language) implementations use a runtime stack to hold parameters, local variables, and other information
  - Stack needed to implement recursion
- Function call mechanism will impact performance
  - A little time needed for function call/return (overhead)
  - Memory also required for stack; implications on recursion
  - In a multi-threaded program, each thread has its own runtime stack, introducing some complexity

# Von Neumann Machine Model



- Most Input/Output (I/O) devices are EXTREMELY slow compared to CPU
  - Disk: read/write ~1,000,000 times slower than CPU register
  - Keyboard/display: much slower than disk!
  - Doesn't make sense for CPU to remain idle waiting for I/O operations to complete
- I/O devices operate concurrent with CPU
  - How does the CPU know when an I/O operation has completed?

# Device Registers




CPU communicates with I/O devices through device registers

- Data (transferred to/from device)
- Control: CPU writes into this register to (for example) initiate an I/O operation and specify parameters (e.g., what disk locations)
- Status: CPU reads this register to find out the status of I/O operations (has data been received? I/O errors?)


# Machine Instructions to Access Device Register

- I/O instructions
  - Machine instructions defined specifically to read/write device registers
  - Separate “address space” for I/O devices
- Memory-Mapped I/O (more common)
  - Each device register assigned to a specific memory address
  - Existing instructions (e.g., LDR, STR) used to read/write device registers

Example CPU I/O instruction




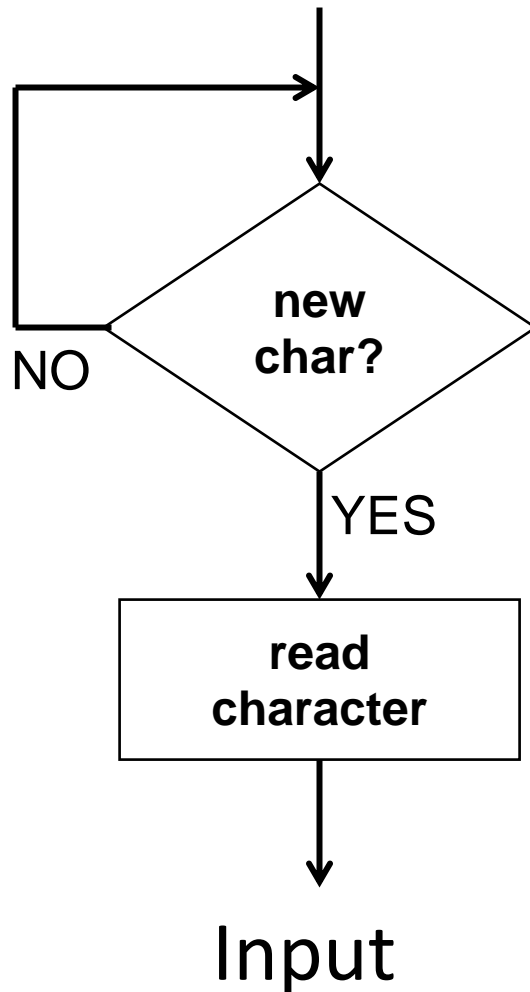
Example:  
Memory mapped I/O



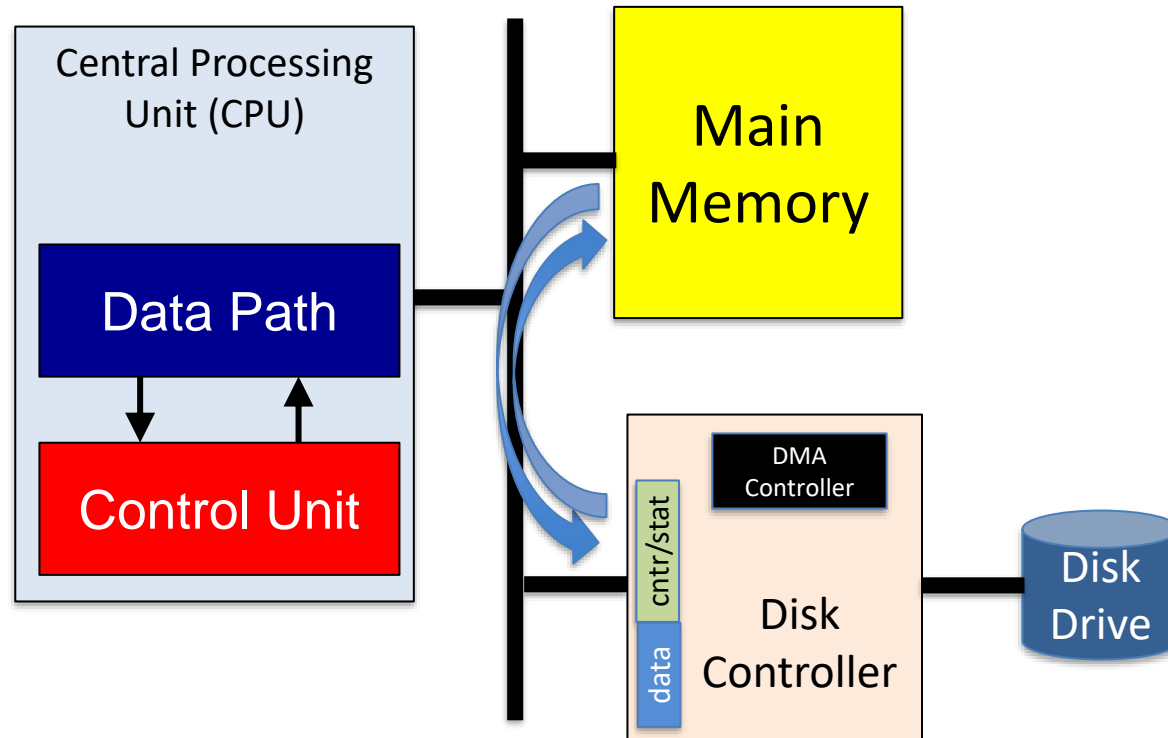


# Programmed I/O: Polling

Low level I/O software usually implemented as functions **within the operating system** (device drivers)



# Direct Memory Access (DMA)



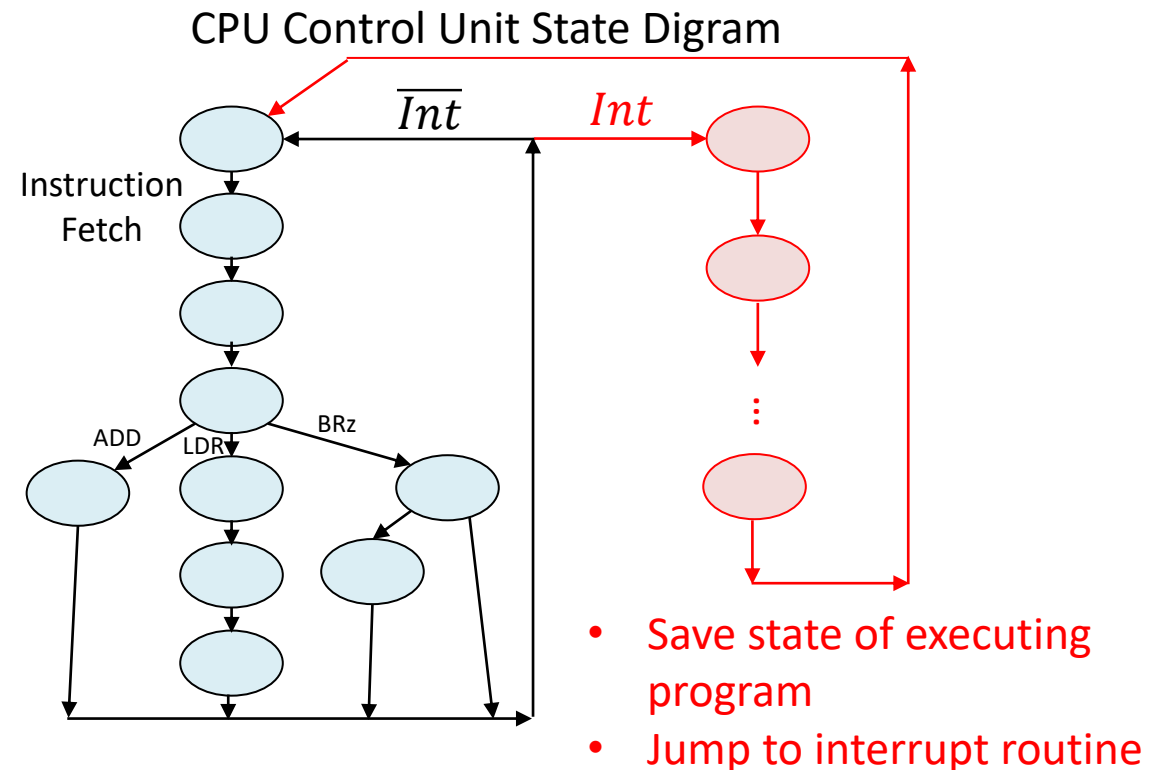
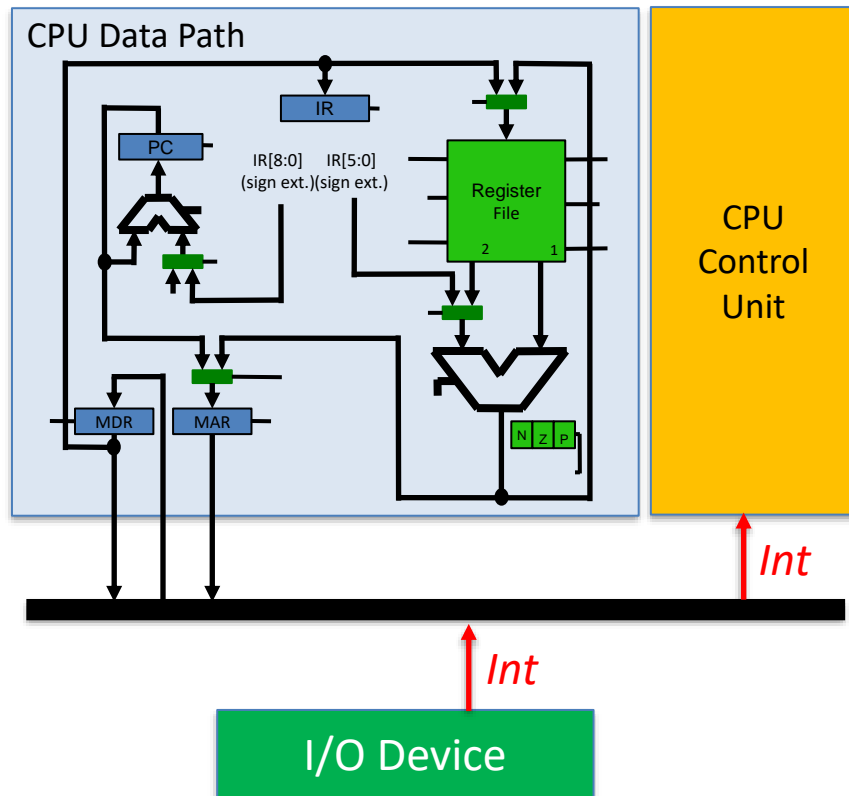
- High-speed devices: programmed I/O uses too much CPU time
  - Disk drives (rotating, solid state): read/write ~4K bytes at a time
  - High-speed network interface
  - Graphics (may use main memory or memory within graphics card)
- Programmed I/O used to start/finish I/O operations
- Data: Device directly accesses main memory **independent of the CPU**
- Leads to an alternative to polling – next topic!

# Polling vs. Interrupts

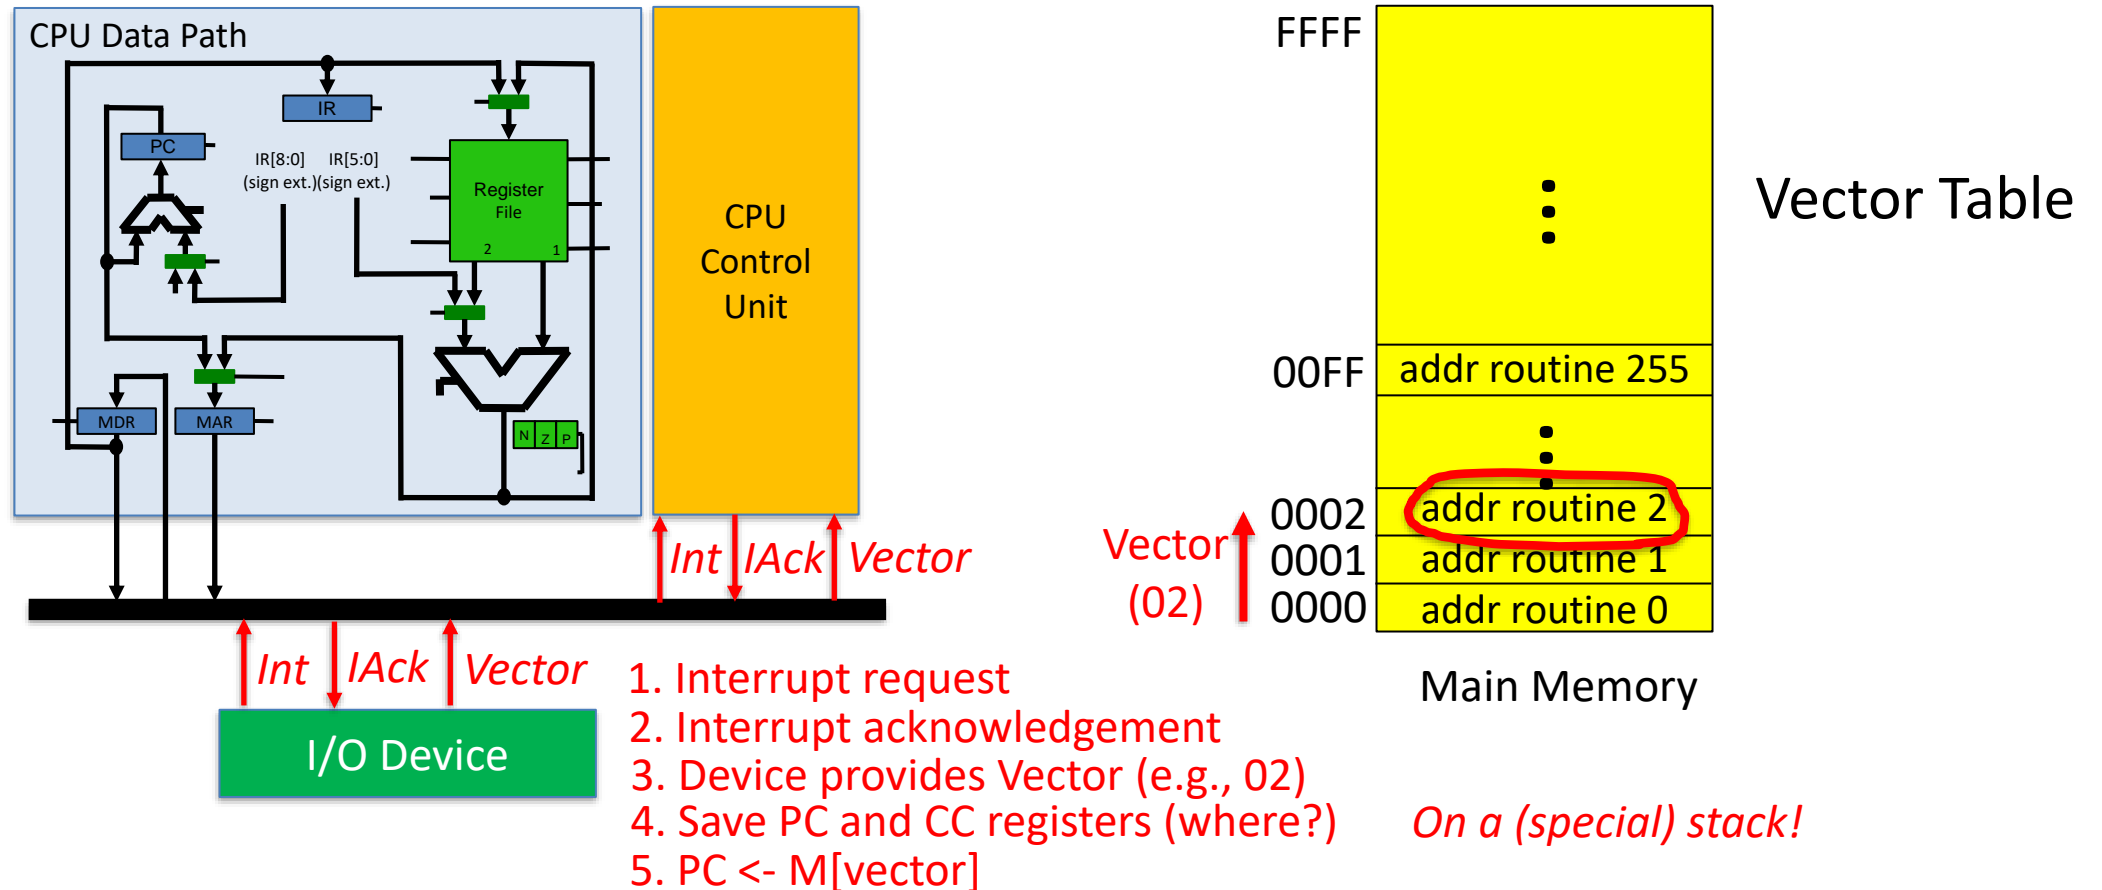
- Polling
  - CPU keeps checking status register until new data arrives OR device ready for next data
  - “Are we there yet? Are we there yet? Are we there yet?”
- Interrupts
  - Device sends a special signal to CPU when new data arrives OR device ready for next data
  - CPU can be performing other tasks instead of polling device.
  - “Wake me when we get there.”
- To implement an interrupt mechanism, we need:
  - A way for the I/O device to **signal** the CPU that an interesting event has occurred
  - A way for the CPU to **test** whether the interrupt signal is set
  - A way to identify which device generated the interrupt

# Implementing Interrupts: Hardware

- Add Interrupt Status Signal (*Int*)
  - Asserted by I/O device to trigger an interrupt
- Hardware checks *Int* before each instruction fetch
- If *Int* is asserted
  - Execute a sequence of steps to process interrupt
  - Resume normal instruction execution



# Vectored Interrupts



- Device provides an interrupt vector to identify itself (e.g., 8 bits)
  - Each device assigned a unique vector
- Vector used as index into table that holds the *starting addresses* of interrupt service routines
  - Location of table often fixed (defined by hardware) [e.g., 0000 to 00FF]

# Interrupts: Saving Processor State

- The program being interrupted does not know it is being interrupted!
- Interrupt mechanism must save enough state from the currently executing program so its execution can be resumed later
  - LC-3: PC and CC
  - Other state (e.g., registers R0-R7 saved by the interrupt handler software if the interrupt handler uses them)
- Return to the interrupted program
  - Restore state of the interrupted program (PC, CC)
    - $PC \leftarrow \text{address of the interrupted instruction} + 1$
  - Return often implemented with a special machine instruction
    - RTI: Return from Interrupt (LC-3; restores CC register)

# Other Aspects

- An interrupt handler might itself be interrupted by another device (nested interrupts)
  - Use “system” runtime stack (separate from the “user” runtime stack discussed earlier)
- An interrupt handler might not want to be interrupted
  - CPUs have the ability to “turn off” interrupts
  - Interrupt priorities often used to determine which interrupts can interrupt which interrupt handlers

# When to Use Interrupts

- When timing of external events is uncertain
  - Example: incoming packet from network
- When device operation takes a long time
  - Example: start a disk transfer, disk interrupts when transfer is finished
  - processor can do something else in the meantime
- When event is rare but critical
  - Example: machine is losing power!



# Traps

- A **Trap** is essentially the same as an interrupt, except it is *generated internally by the CPU* rather than by an external device
- User program requests the OS to perform a task on behalf of the program; OS takes control, handles the request, and returns control to the user program
- Examples
  - System call, e.g., to initiate I/O operations
    - LC-3: TRAP instruction defined for this purpose
    - Like a function call, except the processor changes to “system mode” which can access device registers, and perform other functions user code cannot
    - Execution returns to user program after system call completed
  - Program execution errors
    - Divide by zero, invalid memory reference, etc.
    - Execution jumps to operating system code to “kill” program

# Summary: I/O

- Key concepts (vocabulary)
  - Device registers
  - Memory-mapped I/O
  - Polling
  - Direct memory access (DMA)
  - Interrupts
  - Interrupt vector
  - Trap
- Your program routinely gets interrupted at unpredictable times, affecting the timing of its execution
  - Race conditions in parallel programs
- I/O, interrupt handlers, traps, programmed by system software developers